


Lab 4

Zara Waheed

25th Feb 2022

Question 7.3

```
X <- -2:2
Y <- 1 + 1*X - 2*((X - 1)^2)*I(X >= 1)
plot(X, Y, type = "l")
```



Question 7.9

a)

```
set.seed(1)
fit7.9a <- lm(nox ~ poly(dis, 3), data = Boston)
summary(fit7.9a)
```


```
##
## Call:
## lm(formula = nox ~ poly(dis, 3), data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.121130 -0.040619 -0.009738  0.023385  0.194904
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
##
```

```
## (Intercept)    0.554695    0.002759 201.021 < 2e-16 ***
## poly(dis, 3)1 -2.003096    0.062071 -32.271 < 2e-16 ***
## poly(dis, 3)2  0.856330    0.062071  13.796 < 2e-16 ***
## poly(dis, 3)3 -0.318049    0.062071  -5.124 4.27e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.06207 on 502 degrees of freedom
## Multiple R-squared:  0.7148, Adjusted R-squared:  0.7131
## F-statistic: 419.3 on 3 and 502 DF,  p-value: < 2.2e-16

lim <- range(Boston$dis)
grid <- seq(lim[1], lim[2])
pred <- predict(fit7.9a, list(dis = grid), se = TRUE)


se <- cbind(pred$fit + 2*pred$se.fit, pred$fit - 2*pred$se.fit)

plot(Boston$dis, Boston$nox, col = "red")
lines(grid, pred$fit, col = "blue", lwd = 3)
matlines(grid, se, lwd = 3, col = "green", lty = 3)
```




b)

```
set.seed(1)
rss <- rep(NA, 10)
for (i in 1:10){
  fit <- lm(nox ~ poly(dis, i), data = Boston)
  rss[i] <- sum(fit$residuals^2)
}
plot(1:10, rss, type = "l", lwd = 3)
points(which.min(rss), rss[which.min(rss)], col='blue', pch=10)
```



c)

```
errors <- rep(NA, 10)
for (i in 1:10) {
  fit <- glm(nox ~ poly(dis, i), data = Boston)
  errors[i] <- cv.glm(Boston, fit, K = 10)$delta[1]
}
plot(1:10, errors, type = "l", lwd = 3)
points(which.min(errors), errors[which.min(errors)], col='blue', pch=15)
```



d)

```
summary(Boston$dis)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	1.130	2.100	3.207	3.795	5.188	12.127

```

fit7.9d <- lm(nox ~ bs(dis, df = 4), Boston)

summary(fit7.9d)

##
## Call:
## lm(formula = nox ~ bs(dis, df = 4), data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.124622 -0.039259 -0.008514  0.020850  0.193891
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.73447    0.01460  50.306 < 2e-16 ***
## bs(dis, df = 4)1 -0.05810    0.02186  -2.658  0.00812 **
## bs(dis, df = 4)2 -0.46356    0.02366 -19.596 < 2e-16 ***
## bs(dis, df = 4)3 -0.19979    0.04311  -4.634  4.58e-06 ***
## bs(dis, df = 4)4 -0.38881    0.04551  -8.544 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.06195 on 501 degrees of freedom
## Multiple R-squared:  0.7164, Adjusted R-squared:  0.7142
## F-statistic: 316.5 on 4 and 501 DF,  p-value: < 2.2e-16


attr(bs(Boston$dis, df = 4), "knots")

##      50%
## 3.20745

x <- seq(min(Boston$dis), max(Boston$dis))
y <- predict(fit7.9d, data.frame(dis = x))

plot(Boston$dis, Boston$nox, col = "blue")
lines(x, y, lwd = 2)

```




At 4 degrees of

freedom, we get the knot at 3.207

e)

```
df_vs_rss <- c()
for (i in 3:20) {
  fit <- lm(nox ~ bs(dis, df = i), data = Boston)
  pred <- predict(fit, data.frame(dis = x))
  df_vs_rss[i] <- sum(fit$residuals^2)
}
plot(1:20, df_vs_rss, xlab = "df", ylab = "RSS", type = "l", lwd = 2)
points(which.min(df_vs_rss), df_vs_rss[which.min(df_vs_rss)], col='red', pch=15)
```




freedom gives us the lowest RSS value.

14 degrees of

f)

```
set.seed(100)
cv <- rep(NA, 20)
for (i in 3:20) {
  fit <- glm(nox ~ bs(dis, df = i), data = Boston)
  cv[i] <- cv.glm(Boston, fit, K = 10)$delta[1]
}
plot(3:20, cv[3:20], xlab = "df", ylab = "MSE", type = "l")
points(which.min(cv), cv[which.min(cv)], col = "blue", pch = 15)
```



14 degrees of freedom gives us the lowest MSE value.

Question 7.10

a)

```
data("College")

# Create test and train datasets
set.seed(100)
train_s <- sample(1:nrow(College), 500)
train <- College[train_s,]
test <- College[-train_s,]

fit7.10a <- regsubsets(Outstate ~ ., train, nvmax = ncol(College)-1, method = "forward")

# FSS
fss_summary <- summary(fit7.10a)
par(mfrow = c(1, 3))
plot(fss_summary$cp, xlab = "Variables", ylab = "CP", type = "l")
min.cp <- min(fss_summary$cp)
std.cp <- sd(fss_summary$cp)
abline(h = min.cp + 0.2 * std.cp, col = "blue", lty = 2)
abline(h = min.cp - 0.2 * std.cp, col = "blue", lty = 2)


# BIC
plot(fss_summary$bic, xlab = "Variables", ylab = "BIC", type='l')
min.bic <- min(fss_summary$bic)
std.bic <- sd(fss_summary$bic)
abline(h = min.bic + 0.2 * std.bic, col = "red", lty = 2)
abline(h = min.bic - 0.2 * std.bic, col = "red", lty = 2)

# Adjusted R^2
```

```
plot(fss_summary$adjr2, xlab = "Variables", ylab = "AR^2", type = "l", ylim = c(0.4, 0.84))
max.ar2 <- max(fss_summary$ar2)
```

```
## Warning in max(fss_summary$ar2): no non-missing arguments to max; returning -Inf
```

```
sd.ar2 <- sd(fss_summary$ar2)
abline(h = max.ar2 + 0.2 * sd.ar2, col = "green", lty = 2)
abline(h = max.ar2 - 0.2 * sd.ar2, col = "green", lty = 2)
```



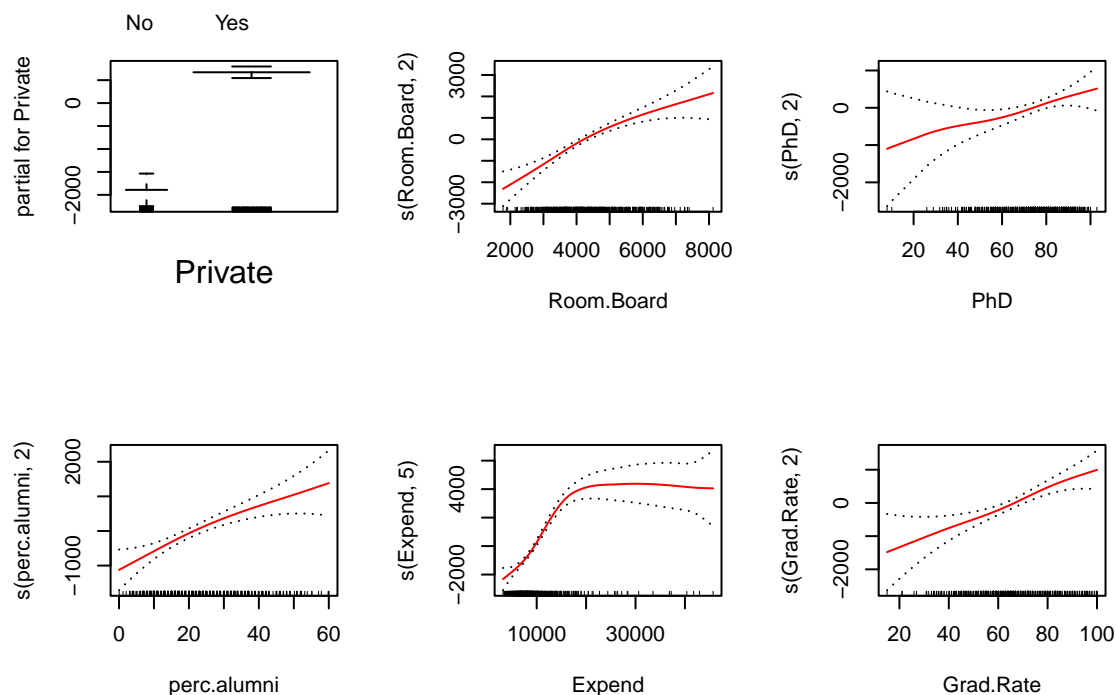
metrics do not seem to improve much after 6 predictors.

b)

```
fit7.10b <- gam(Outstate ~ Private + s(Room.Board,2) + s(PhD,2) + s(perc.alumni,2) + s(Expend,5) + s(Gr
```

```
par(mfrow = c(2,3))
```

```
plot(fit7.10b, se = TRUE, col = "red")
```



d)

```
summary(fit7.10a)
```

```
## Subset selection object
## Call: regsubsets.formula(Outstate ~ ., train, nvmax = ncol(College) -
##      1, method = "forward")
## 17 Variables (and intercept)
##               Forced in Forced out
## PrivateYes      FALSE      FALSE
## Apps            FALSE      FALSE
## Accept          FALSE      FALSE
## Enroll          FALSE      FALSE
## Top10perc       FALSE      FALSE
## Top25perc       FALSE      FALSE
## F.Undergrad     FALSE      FALSE
## P.Undergrad     FALSE      FALSE
## Room.Board      FALSE      FALSE
## Books           FALSE      FALSE
## Personal        FALSE      FALSE
## PhD             FALSE      FALSE
## Terminal        FALSE      FALSE
## S.F.Ratio       FALSE      FALSE
## perc.alumni     FALSE      FALSE
## Expend          FALSE      FALSE
## Grad.Rate       FALSE      FALSE
## 1 subsets of each size up to 17
## Selection Algorithm: forward
##               PrivateYes Apps Accept Enroll Top10perc Top25perc F.Undergrad
## 1  ( 1 )  " "          " " " "  " "  " "          " "
## 2  ( 1 )  "*"        " " " "  " "  " "          " "
```


## 3	(1)	"*	" "	" "	" "	" "	" "	" "	
## 4	(1)	"*	" "	" "	" "	" "	" "	" "	
## 5	(1)	"*	" "	" "	" "	" "	" "	" "	
## 6	(1)	"*	" "	" "	" "	" "	" "	" "	
## 7	(1)	"*	" "	" "	" "	" "	" "	" "	
## 8	(1)	"*	" "	"*	" "	" "	" "	" "	
## 9	(1)	"*	" "	"*	"*	" "	" "	" "	
## 10	(1)	"*	"*	"*	"*	" "	" "	" "	
## 11	(1)	"*	"*	"*	"*	" "	" "	" "	
## 12	(1)	"*	"*	"*	"*	" "	" "	" "	
## 13	(1)	"*	"*	"*	"*	" "	" "	" "	
## 14	(1)	"*	"*	"*	"*	"*	" "	" "	
## 15	(1)	"*	"*	"*	"*	"*	" "	" "	
## 16	(1)	"*	"*	"*	"*	"*	"*	"*	
## 17	(1)	"*	"*	"*	"*	"*	"*	"*	
##			P.Undergrad	Room.Board	Books	Personal	PhD	Terminal	S.F.Ratio
## 1	(1)	" "	" "	" "	" "	" "	" "	" "	" "
## 2	(1)	" "	" "	" "	" "	" "	" "	" "	" "
## 3	(1)	" "	"*	" "	" "	" "	" "	" "	" "
## 4	(1)	" "	"*	" "	" "	" "	" "	" "	" "
## 5	(1)	" "	"*	" "	" "	" "	"*	" "	" "
## 6	(1)	" "	"*	" "	" "	" "	"*	" "	" "
## 7	(1)	" "	"*	" "	"*	" "	"*	" "	" "
## 8	(1)	" "	"*	" "	"*	" "	"*	" "	" "
## 9	(1)	" "	"*	" "	"*	" "	"*	" "	" "
## 10	(1)	" "	"*	" "	"*	" "	"*	" "	" "
## 11	(1)	" "	"*	" "	"*	" "	"*	" "	" "
## 12	(1)	" "	"*	" "	"*	" "	"*	" "	"*
## 13	(1)	" "	"*	"*	"*	" "	"*	" "	"*
## 14	(1)	" "	"*	"*	"*	" "	"*	" "	"*
## 15	(1)	" "	"*	"*	"*	"*	"*	" "	"*
## 16	(1)	" "	"*	"*	"*	"*	"*	" "	"*
## 17	(1)	"*	"*	"*	"*	"*	"*	" "	"*
##			perc.alumni	Expend	Grad.Rate				
## 1	(1)	" "	"*	" "	" "				
## 2	(1)	" "	"*	" "	" "				
## 3	(1)	" "	"*	" "	" "				
## 4	(1)	"*	"*	" "	" "				
## 5	(1)	"*	"*	" "	" "				
## 6	(1)	"*	"*	"*	" "				
## 7	(1)	"*	"*	"*	" "				
## 8	(1)	"*	"*	"*	" "				
## 9	(1)	"*	"*	"*	" "				
## 10	(1)	"*	"*	"*	" "				
## 11	(1)	"*	"*	"*	" "				
## 12	(1)	"*	"*	"*	" "				
## 13	(1)	"*	"*	"*	" "				
## 14	(1)	"*	"*	"*	" "				
## 15	(1)	"*	"*	"*	" "				
## 16	(1)	"*	"*	"*	" "				
## 17	(1)	"*	"*	"*	" "				

The relationship seems to be non-linear

Question 7.11

a)

```
set.seed(100)
y <- rnorm(100)
x1 <- rnorm(100)
x2 <- rnorm(100)
```

b)

```
b1 <- 1
```

c)

```
a <- y - b1*x1
b2 <- lm(a~x2)$coef[2]
```


d)

```
a <- y - b2*x2
b1 <- lm(a~x1)$coef[2]
```

e)

```
iterations <- 10
df <- data.frame(0.0, 0.27, 0.0)
names(df) <- c('b0', 'b1', 'b2')
for (i in 1:iterations) {
  b1 <- df[nrow(df), 2]
  a <- y - b1 * x1
  b2 <- lm(a ~ x2)$coef[2]
  a <- y - b2 * x2
  b1 <- lm(a ~ x1)$coef[2]
  b0 <- lm(a ~ x1)$coef[1]
  b0
  b1
  b2
  df[nrow(df) + 1,] <- list(b0, b1, b2)
}


plot(df$b0, col = 'red', type = 'l')
lines(df$b1, col = 'blue')
lines(df$b2, col = 'green')
```



f)

```
plot(df$b0, col = 'yellow', type = 'l')
lines(df$b1, col = 'blue')
lines(df$b2, col = 'red')

d <- coef(lm(y ~ x1 + x2))
abline(h = d[1], col = 'orange', lty = 3)
abline(h = d[2], col = 'purple', lty = 3)
abline(h = d[3], col = 'pink', lty = 3)
```



g)

More than 5 iterations were required for a good approximation.